

Arrows and Reagents

“KC” Sivaramakrishnan

Advanced Functional Programming

March 3rd, 2016

Arrows

```
module type Arrow =  
sig  
  type ('a,'b) t  
  val arr    : ('a -> 'b) -> ('a,'b) t  
  val (>>>) : ('a,'b) t -> ('b,'c) t -> ('a,'c) t  
  val first : ('a,'b) t -> ('a * 'c, 'b * 'c) t  
end
```

Arrows

```
module type Arrow =  
sig  
  type ('a,'b) t  
  val arr    : ('a -> 'b) -> ('a,'b) t  
  val (>>>) : ('a,'b) t -> ('b,'c) t -> ('a,'c) t  
  val first  : ('a,'b) t -> ('a * 'c, 'b * 'c) t  
end
```

Laws

```
arr f >>> arr g ≡ arr (compose g f)  
(f >>> g) >>> h ≡ f >>> (g >>> h)  
arr id >>> f ≡ f  
... ..
```

Functions as Arrows

- <https://gist.github.com/9eef070c232913121564>

“If we think of a library as defining a domain specific '**language**', whose constructions are represented as **combinators**, then the idea is to implement the language via a combination of a **static analysis** and **an optimised dynamic semantics.**”

John Hughes, “Generalising Monads to Arrows”

“If we think of a library as defining a domain specific '**language**', whose constructions are represented as **combinators**, then the idea is to implement the language via a combination of a **static analysis** and **an optimised dynamic semantics**.”

John Hughes, “Generalising Monads to Arrows”

```
val (>>=) : 'a Monad.t -> ('a -> 'b Monad.t) -> 'b Monad.t
```

```
val (>>>) : ('a, 'b) Arrow.t -> ('b, 'c) Arrow.t -> ('a, 'c) Arrow.t
```

Functions with *cost* as Arrows

- <https://gist.github.com/66fcc8c01b563282ef42>
- <https://gist.github.com/644fbe3d36f90d98faa1>

Reagents

- DSL for *expressing* and *composing* fine-grained concurrency libraries
- Aaron Turon, “Reagents: expressing and composing fine-grained concurrency”, PLDI 2012
- Based on Arrows
 - Enable dynamic optimisations
- Built on *k-compare-and-swap* abstraction

Compare-and-swap (CAS)


```
module CAS : sig
  val cas : 'a ref -> expect:'a -> update:'a -> bool
end = struct
  (* atomically... *)
  let cas r ~expect ~update =
    if !r = expect then
      (r := update; true)
    else false
  end
```

Compare-and-swap (CAS)

```
module CAS : sig
  val cas : 'a ref -> expect:'a -> update:'a -> bool
end = struct
  (* atomically... *)
  let cas r ~expect ~update =
    if !r = expect then
      (r := update; true)
    else false
  end
```

- Implemented *atomically* by processors
 - x86: CMPXCHG and friends
 - arm: LDREX, STREX, etc.
 - ppc: lwarx, stwcx, etc.

CAS: cost versus contention



java.util.concurrent

Synchronization

Reentrant locks

Semaphores

R/W locks

Reentrant R/W locks

Condition variables

Countdown latches

Cyclic barriers

Phasers

Exchangers

Data structures

Queues

Nonblocking

Blocking (array & list)

Synchronous

Priority, nonblocking

Priority, blocking

Dequeues

Sets

Maps (hash & skiplist)

java.util.concurrent

Synchronization

Reentrant locks

Semaphores

R/W locks

Reentrant R/W locks

Condition variables

Countdown

Cyclic

PL

Exchangers

Data

Ordered

Blocking

Blocking (array & list)

Synchronous

Priority, nonblocking

Priority, blocking

Dequeues

Sets

Maps (hash & skiplist)

Not Composable

```


module type TREIBER_STACK = sig
  type 'a t
  val push : 'a t -> 'a -> unit
  ...
end

module Treiber_stack : TREIBER_STACK =
struct
  type 'a t = 'a list ref


  let rec push s t =
    let cur = !s in
    if CAS.cas s cur (t::cur) then ()
    else (backoff (); push s t)
  end
end


```


Head





Head











```

module type TREIBER_STACK = sig
  type 'a t
  val push      : 'a t -> 'a -> unit
  val try_pop   : 'a t -> 'a option
end

module Treiber_stack : TREIBER_STACK =
struct
  type 'a t = 'a list ref

  let rec push s t = ...

  let rec try_pop s =
    match !s with
    | [] -> None
    | (x::xs) as cur ->
        if CAS.cas s cur xs then Some x
        else (backoff (); try_pop s)
end

```

The Problem:

Concurrency libraries are
indispensable, but hard to
build and extend

```
let v = Treiber_stack.pop s1 in  
Treiber_stack.push s2 v
```

is not *atomic*

The Proposal:


Scalable concurrent algorithms
can be **built** and **extended** using
abstraction and **composition**

```
Treiber_stack.pop s1 >>> Treiber_stack.push s2
```

is ***atomic***


Design

Lambda: the ultimate abstraction




`val f : 'a -> 'b`

Lambda: the ultimate abstraction




`val f : 'a -> 'b`




`val g : 'b -> 'c`

Lambda: the ultimate abstraction



$(\text{compose } g \ f): 'a \rightarrow 'c$


Lambda abstraction:



Lambda abstraction:



Reagent abstraction:



('a, 'b) Reagent . t

Reagent combinators

```
module type Reagents = sig
  type ('a,'b) t
  val never      : ('a,'b) t
  val constant   : 'a -> ('b,'a) t
  val (>>>)      : ('a,'b) t -> ('b,'c) t -> ('a,'c) t

  module Ref : Ref.S with type ('a,'b) reagent = ('a,'b) t
  module Channel : Channel.S with type ('a,'b) reagent = ('a,'b) t

  val run      : ('a,'b) t -> 'a -> 'b
  ...
end
```

```
module type Channel = sig
  type ('a,'b) endpoint
  type ('a,'b) reagent

  val mk_chan : unit -> ('a,'b) endpoint * ('b,'a) endpoint
  val swap    : ('a,'b) endpoint -> ('a,'b) reagent
end
```


```

module type Channel = sig
  type ('a,'b) endpoint
  type ('a,'b) reagent

  val mk_chan : unit -> ('a,'b) endpoint * ('b,'a) endpoint
  val swap    : ('a,'b) endpoint -> ('a,'b) reagent
end

```

c: ('a,'b) endpoint





```

module type Channel = sig
  type ('a,'b) endpoint
  type ('a,'b) reagent


  val mk_chan : unit -> ('a,'b) endpoint * ('b,'a) endpoint
  val swap    : ('a,'b) endpoint -> ('a,'b) reagent
end

```


c: ('a,'b) endpoint



c: ('a, 'b) endpoint




Message passing




```
type 'a ref  
val upd : 'a ref  
-> f:('a -> 'b -> ('a * 'c) option)  
-> ('b, 'c) Reagent.t
```


Message passing




```
type 'a ref  
val upd : 'a ref  
-> f:('a -> 'b -> ('a * 'c) option)  
-> ('b, 'c) Reagent.t
```




Message passing





Shared state




Message passing





Shared state




Message passing




Shared state




Message passing




Shared state




Disjunction





Message passing




Shared state




Disjunction





Message passing




Shared state




Disjunction




Message passing




Shared state



Disjunction



Conjunction



```

module type TREIBER_STACK = sig
  type 'a t
  val create    : unit -> 'a t
  val push      : 'a t -> ('a, unit) Reagent.t
  val pop       : 'a t -> (unit, 'a) Reagent.t
  val try_pop   : 'a t -> (unit, 'a option) Reagent.t
end

module Treiber_stack : TREIBER_STACK = struct
  type 'a t = 'a list Ref.ref

  let create () = Ref.mk_ref []

  let push r x = Ref.upd r (fun xs x -> Some (x::xs, ()))

  let try_pop r = Ref.upd r (fun l () ->
    match l with
    | [] -> Some ([], None)
    | x::xs -> Some (xs, Some x))

  let pop r = Ref.upd r (fun l () ->
    match l with
    | [] -> None
    | x::xs -> Some (xs, x))
end

```

Composability

Transfer elements atomically

```
Treiber_stack.pop s1 >>> Treiber_stack.push s2
```

Composability

Transfer elements atomically

```
Treiber_stack.pop s1 >>> Treiber_stack.push s2
```

Consume elements atomically

```
Treiber_stack.pop s1 <*> Treiber_stack.pop s2
```

Composability

Transfer elements atomically


```
Treiber_stack.pop s1 >>> Treiber_stack.push s2
```

Consume elements atomically

```
Treiber_stack.pop s1 <*> Treiber_stack.pop s2
```

Consume elements from either

```
Treiber_stack.pop s1 <+> Treiber_stack.pop s2
```



```
type fork =  
  {drop : (unit,unit) endpoint;  
   take : (unit,unit) endpoint}
```


```
let mk_fork () =  
  let drop, take = mk_chan () in  
  {drop; take}
```

```
let drop f = swap f.drop  
let take f = swap f.take
```


```
let init forks =  
  List.iter (fun fork ->  
    Thread.spawn @@ run (drop fork)) forks
```


```
let eat l_fork r_fork =  
  run (take l_fork <*> take r_fork) ();  
  (* ...  
   * eat  
   * ... *)  
  run (drop l_fork) ();  
  run (drop r_fork) ()
```


Implementation











Accumulate CASes


Attempt k-CAS





Accumulate CASes


Attempt k-CAS

















$$P \ \& \ P = P$$

$$T \ \& \ T = T$$

$$P \ \& \ T = T$$

$$T \ \& \ P = T$$

Trouble with k-CAS

Trouble with k-CAS

- Most processors do not support k-CAS

Trouble with k-CAS

- Most processors do not support k-CAS
- Implemented as a multi-phase protocol
 1. Sort refs
 2. Lock refs in order (CAS); rollback if conflicts.
 3. Commit refs

Trouble with k-CAS

- Most processors do not support k-CAS
- Implemented as a multi-phase protocol
 1. Sort refs
 2. Lock refs in order (CAS); rollback if conflicts.
 3. Commit refs
- Additional book-keeping required
 - CAS list, messages to be consumed, post-commit actions, etc.

Trouble with k-CAS

- Most processors do not support k-CAS
- Implemented as a multi-phase protocol
 1. Sort refs
 2. Lock refs in order (CAS); rollback if conflicts.
 3. Commit refs
- Additional book-keeping required
 - CAS list, messages to be consumed, post-commit actions, etc.
- Common case is *just a single CAS*
 - *Identify and optimise with Arrows*

Reagent type

```
type 'a result = Block | Retry | Done of 'a
```

```
type ('a,'b) t =  
  { try_react : 'a -> Reaction.t -> 'b Offer.t option -> 'b result;  
    compose : 'r. ('b,'r) t -> ('a,'r) t;  
    always_commits : bool;  
    may_sync : bool }
```

Reagent type

permanent failure




```
type 'a result = Block | Retry | Done of 'a
```

```
type ('a, 'b) t =  
  { try_react : 'a -> Reaction.t -> 'b Offer.t option -> 'b result;  
    compose : 'r. ('b, 'r) t -> ('a, 'r) t;  
    always_commits : bool;  
    may_sync : bool }
```

Reagent type

permanent failure transient failure



```
type 'a result = Block | Retry | Done of 'a
```

```
type ('a,'b) t =  
  { try_react : 'a -> Reaction.t -> 'b Offer.t option -> 'b result;  
    compose : 'r. ('b,'r) t -> ('a,'r) t;  
    always_commits : bool;  
    may_sync : bool }
```

Reagent type

permanent failure transient failure CAS set

```
type 'a result = Block | Retry | Done of 'a
```

type ('a, 'b) t =
 { try_react : 'a -> Reaction.t -> 'b Offer.t option -> 'b result;
 compose : 'r. ('b, 'r) t -> ('a, 'r) t;
 always_commits : bool;
 may_sync : bool }

The diagram illustrates the relationship between the Reagent type and its components. It shows three arrows: one from the 'Block' constructor to the label 'permanent failure', one from the 'Retry' constructor to the label 'transient failure', and one from the 'Reaction.t' type in the 'try_react' field to the label 'CAS set'.

Reagent type

permanent failure transient failure CAS set

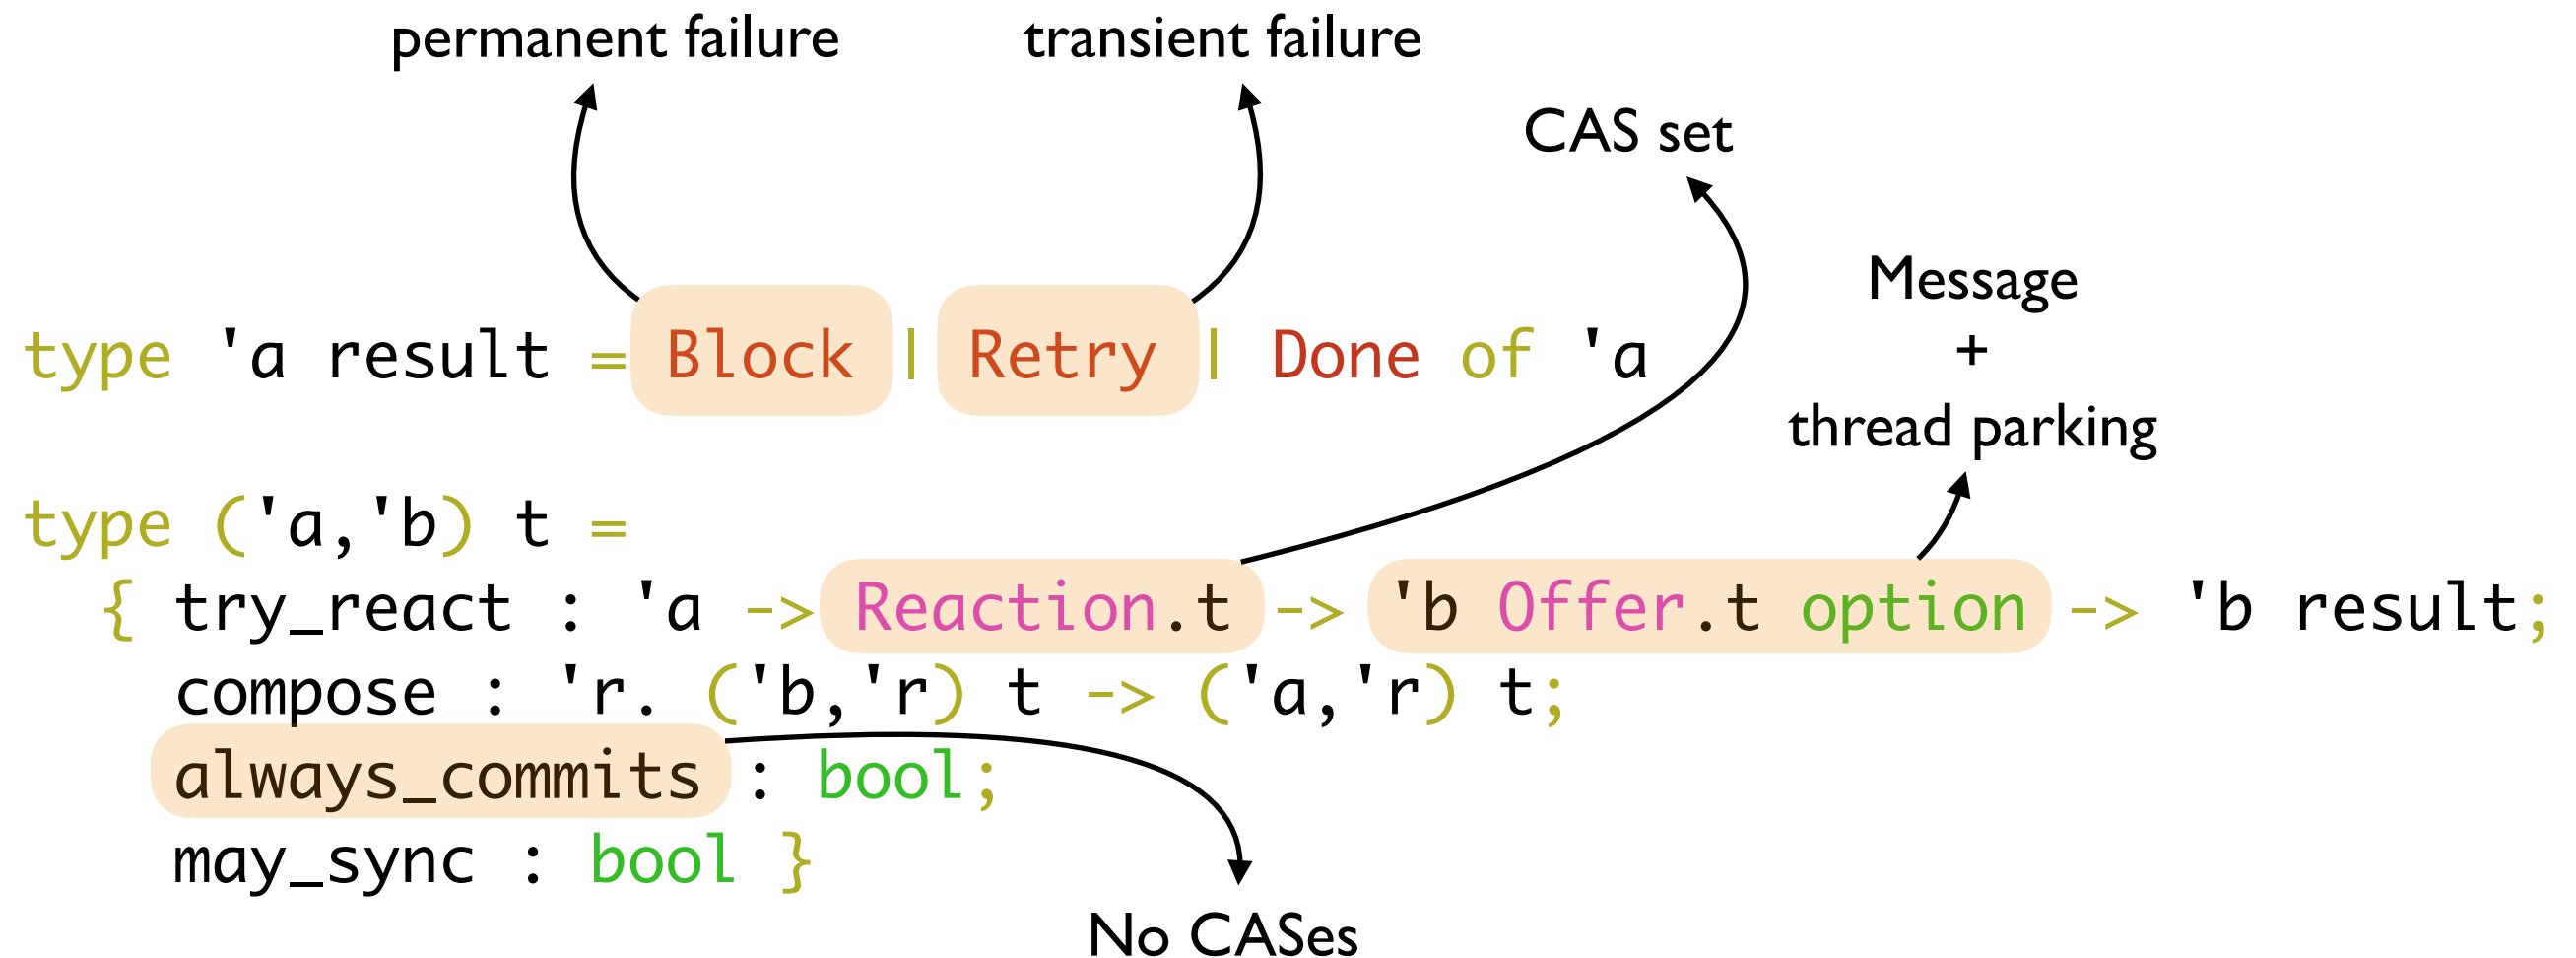
```
type 'a result = Block | Retry | Done of 'a
```

Message
+
thread parking

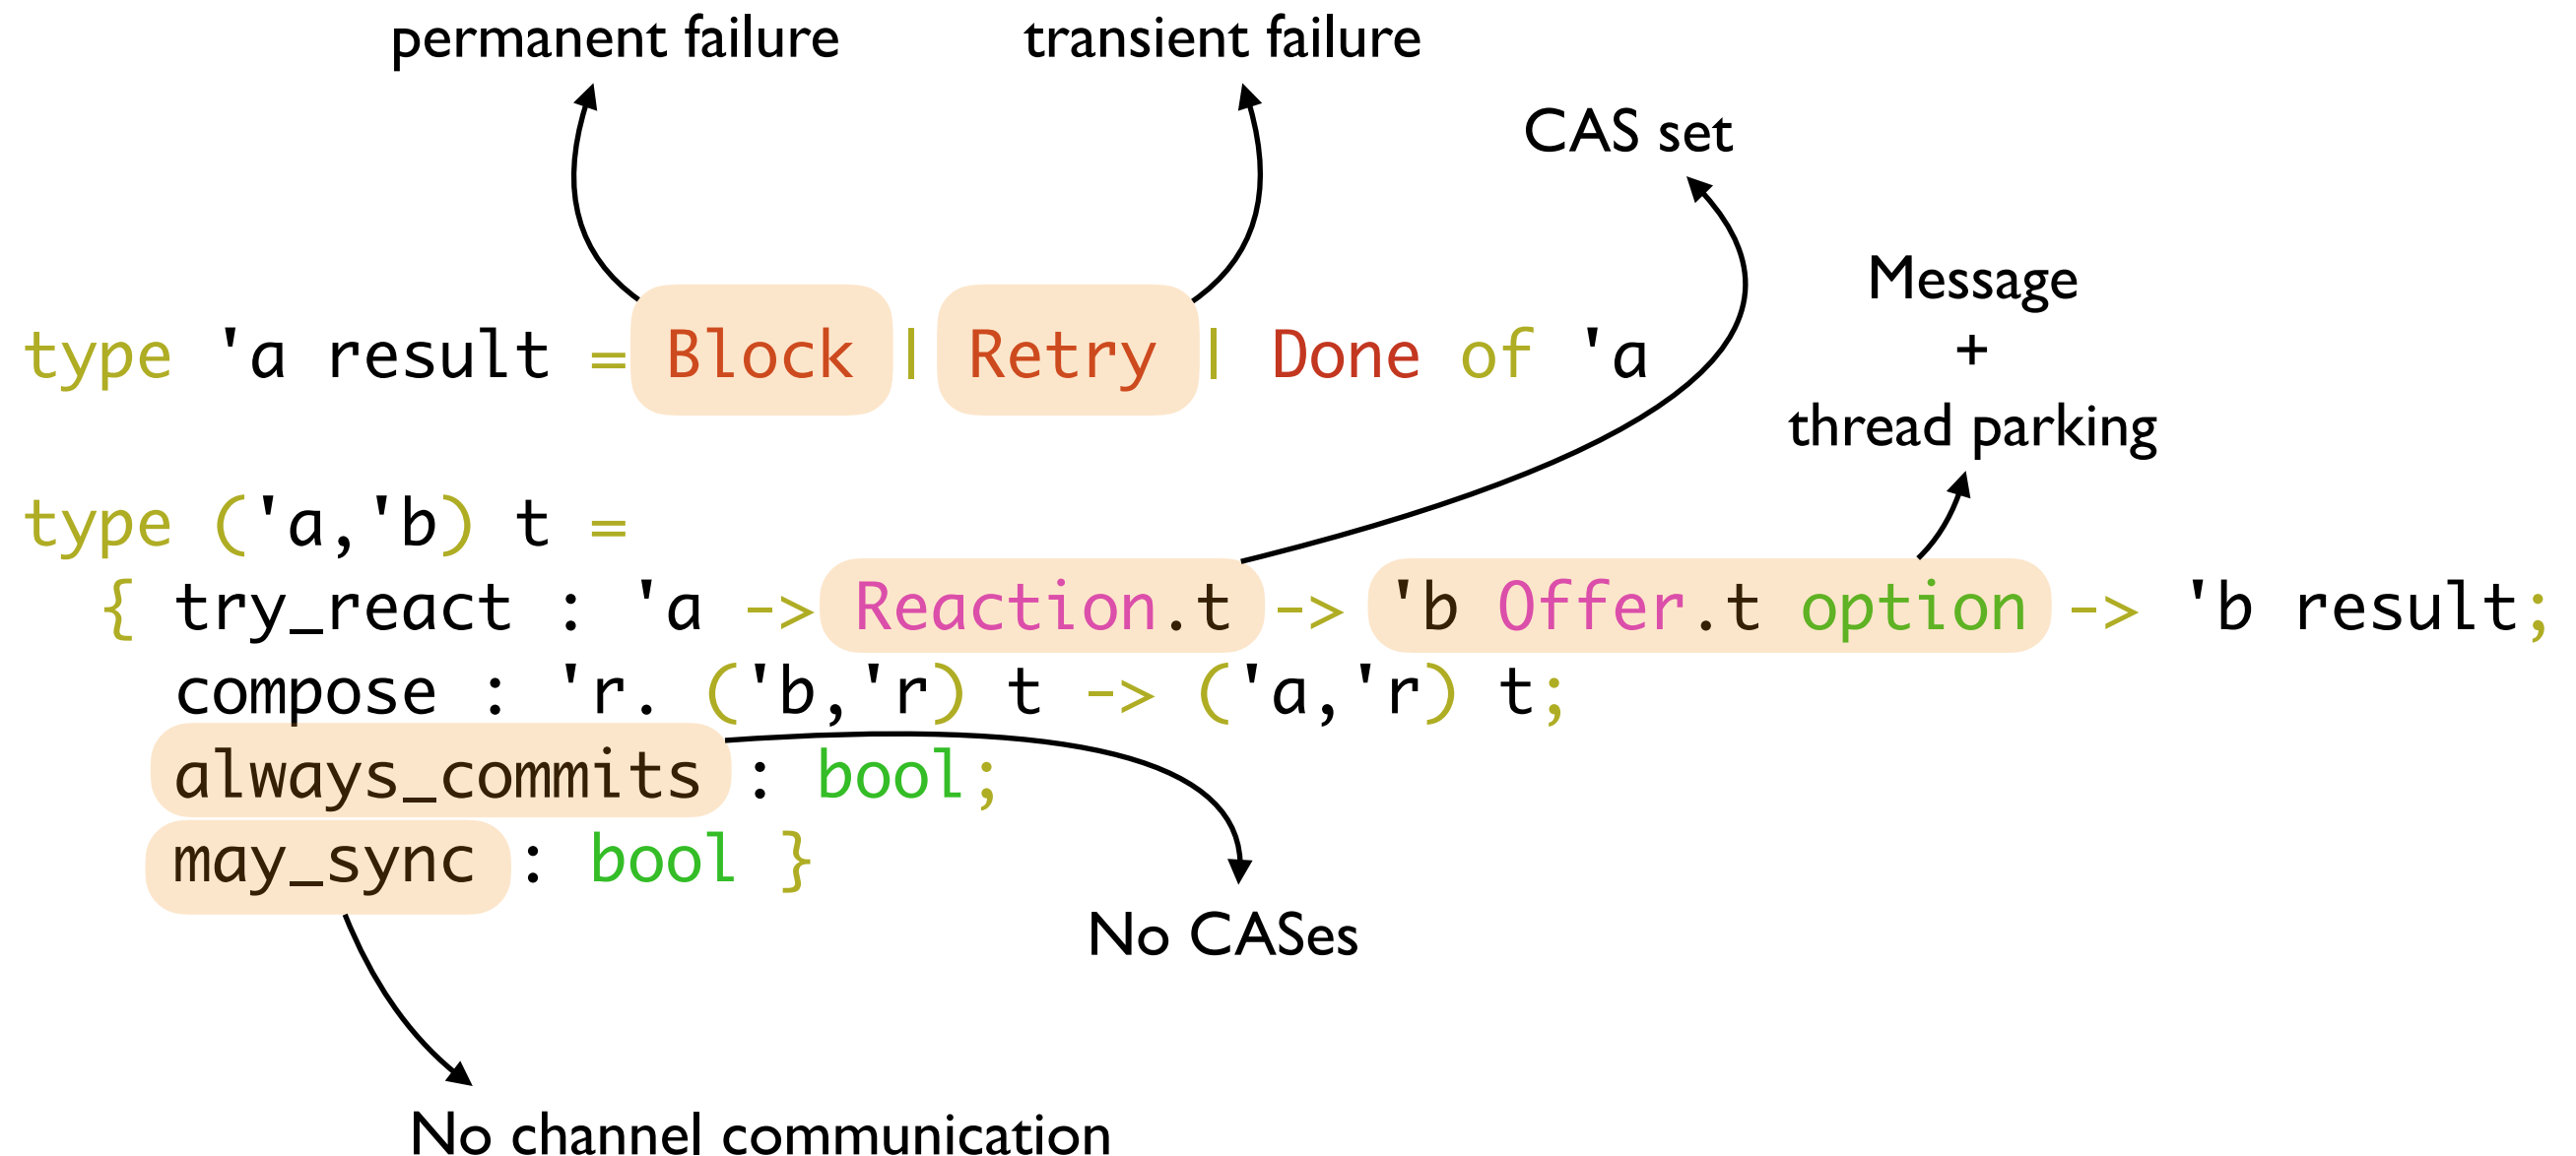
```
type ('a, 'b) t =  
  { try_react : 'a -> Reaction.t -> 'b Offer.t option -> 'b result;  
    compose : 'r. ('b, 'r) t -> ('a, 'r) t;  
    always_commits : bool;  
    may_sync : bool }
```

The diagram illustrates the components of the Reagent type. It shows two type definitions. The first defines 'a result as a union of Block, Retry, and Done of 'a. The second defines ('a, 'b) t as a record with four fields: try_react, compose, always_commits, and may_sync. Annotations with arrows point from specific parts of the code to descriptive text: 'Block' is annotated with 'permanent failure', 'Retry' with 'transient failure', 'Reaction.t' with 'CAS set', and 'Offer.t option' with 'Message + thread parking'.

Reagent type



Reagent type



```
let rec never : 'a 'b. ('a, 'b) t =  
  { try_react = (fun _ _ _ -> Block);  
    may_sync = false;  
    always_commits = false;  
    compose = fun _ -> never }
```

```
let rec never : 'a 'b. ('a,'b) t =  
  { try_react = (fun _ _ _ -> Block);  
    may_sync = false;  
    always_commits = false;  
    compose = fun _ -> never }
```

```
let rec constant : 'a 'b 'r. 'a -> ('a,'r) t -> ('b, 'r) t =  
  fun x k (* continuation *) ->  
    { may_sync = k.may_sync;  
      always_commits = k.always_commits;  
      try_react = (fun _ rx o -> k.try_react x rx o);  
      compose = (fun next -> constant x (k.compose next)) }
```


```
let rec never : 'a 'b. ('a,'b) t =
  { try_react = (fun _ _ _ -> Block);
    may_sync = false;
    always_commits = false;
    compose = fun _ -> never }
```

```
let rec constant : 'a 'b 'r. 'a -> ('a,'r) t -> ('b, 'r) t =
  fun x k (* continuation *) ->
    { may_sync = k.may_sync;
      always_commits = k.always_commits;
      try_react = (fun _ rx o -> k.try_react x rx o);
      compose = (fun next -> constant x (k.compose next)) }
```

```
let rec <+> : 'a 'b 'r. ('a,'b) t -> ('a,'b) t -> ('a,'b) t =
  fun r1 r2 ->
    { always_commits = r1.always_commits && r1.always_commits;
      may_sync = r1.may_sync || r2.may_sync;
      ... }
```

Specialising k-CAS

```
let rec cas r ~expect ~update k =  
  let try_react () rx o =  
    if Reaction.has_no_cas rx &&  
      k.always_commits then  
      if CAS.cas r.data expect update then  
        ( k.try_react () rx o ) (* Will succeed! *)  
      else Retry  
    else  
      (* slow path with bookkeeping *)  
in  
...
```



Optimising Transient Failures

```
let rec without_offer pause r v =  
  match r.try_react v Reaction.empty None with  
  | Done res -> res  
  | Retry ->  
    ( pause ();  
      if r.may_sync  
      then with_offer pause r v  
      else without_offer pause r v)  
  | Block -> with_offer pause r v
```

```
let run r v =  
  let b = Backoff.create () in  
  let pause () = Backoff.once b in  
  without_offer pause r v
```